

▼ SETUP INICIAL - Instalar librerias y configurar Google Colab

```
# Instalar librerias necesarias
!pip install -q scikit-learn pandas numpy matplotlib seaborn

# Importar librerias
import pandas as pd
import numpy as np
import os
import pickle
import warnings
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, classification_report, roc_auc_score,
    roc_curve, auc
)
warnings.filterwarnings('ignore')

# Configurar visualizaciones
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette("husl")

print("[OK] Librerias importadas exitosamente")
```

[OK] Librerias importadas exitosamente

```
# Configurar Google Colab (si aplica)
from google.colab import files
import shutil

print("[OK] Google Colab configurado")
print("\nINSTRUCCIONES PARA COLAB:")
print("1. Si tienes el archivo CSV localmente, ejecuta: files.upload()")
print("2. O copia el CSV a tu Google Drive y usa:")
print("   from google.colab import drive")
print("   drive.mount('/content/drive')")
print("3. Luego actualiza dataset_path a la ubicacion correcta")

[OK] Google Colab configurado

INSTRUCCIONES PARA COLAB:
1. Si tienes el archivo CSV localmente, ejecuta: files.upload()
2. O copia el CSV a tu Google Drive y usa:
   from google.colab import drive
   drive.mount('/content/drive')
3. Luego actualiza dataset_path a la ubicacion correcta
```

▼ FASE 1: SAMPLE - Cargar y verificar datos

Objetivo: Cargar dataset original y validar estructura

```
print("\n" + "="*80)
print("FASE 1: SAMPLE - Cargar Dataset")
print("="*80)
print("Objetivo: Cargar 9,312 muestras del dataset CyberNative DPO")

dataset_path = '/content/cybernative_detector_training.csv'

print(f"\nCargando dataset desde: {dataset_path}")

try:
    df = pd.read_csv(dataset_path)
    print(f"[OK] Dataset cargado exitosamente")
except FileNotFoundError:
    print(f"[ERROR] Archivo no encontrado: {dataset_path}")
```

```

print("Por favor ajusta la ruta dataset_path")
df = None

=====
FASE 1: SAMPLE - Cargar Dataset
=====
Objetivo: Cargar 9,312 muestras del dataset CyberNative DPO

Cargando dataset desde: /content/cybernative_detector_training.csv
[OK] Dataset cargado exitosamente

```

```

# Validar estructura del dataset
print(f"\nTotal registros: {len(df)}")
print(f"Total columnas: {df.shape[1]}")
print(f"\nColumnas: {list(df.columns)}")

# Verificar columnas requeridas
required_cols = ['codigo', 'lenguaje', 'vulnerable']
missing_cols = [col for col in required_cols if col not in df.columns]
if missing_cols:
    print(f"[ERROR] Faltan columnas: {missing_cols}")
else:
    print(f"[OK] Todas las columnas requeridas presentes")

```

```

Total registros: 9,312
Total columnas: 5

Columnas: ['codigo', 'lenguaje', 'tipo_vulnerabilidad', 'pregunta', 'vulnerable']
[OK] Todas las columnas requeridas presentes

```

```

# Analizar distribucion de clases
print("\nDistribucion de clases:")
class_counts = df['vulnerable'].value_counts()
for cls, count in class_counts.items():
    pct = (count / len(df)) * 100
    label = "Vulnerable" if cls == 1 else "Seguro"
    print(f" - {label} ({cls}): {count}, ({pct:.1f}%)"

# Analizar lenguajes
print("\nLenguajes en dataset: {df['lenguaje'].nunique()}")
print(f"Valores nulos totales: {df.isnull().sum().sum()}")
print(f"Duplicados: {df.duplicated().sum()}")

print("\n[OK] FASE 1 COMPLETADA")

```

```

Distribucion de clases:
- Vulnerable (1): 4,656 (50.0%)
- Seguro (0): 4,656 (50.0%)

Lenguajes en dataset: 11
Valores nulos totales: 0
Duplicados: 7

[OK] FASE 1 COMPLETADA

```

▼ FASE 2: EXPLORE - Analisis exploratorio

Objetivo: Entender caracteristicas del dataset

```

print("\n" + "="*80)
print("FASE 2: EXPLORE - Analisis Estadistico")
print("="*80)

# Estadisticas de longitud de codigo
code_lengths = df['codigo'].str.len()
print(f"\nEstadisticas de longitud de codigo:")
print(f" - Promedio: {code_lengths.mean():.0f} caracteres")
print(f" - Mediana: {code_lengths.median():.0f} caracteres")
print(f" - Minimo: {code_lengths.min()} caracteres")
print(f" - Maximo: {code_lengths.max()} caracteres")
print(f" - Desv. Est.: {code_lengths.std():.0f} caracteres")

```

```

=====
FASE 2: EXPLORE - Analisis Estadistico
=====
```

```
Estadisticas de longitud de codigo:
- Promedio: 461 caracteres
- Mediana: 369 caracteres
- Minimo: 55 caracteres
- Maximo: 8925 caracteres
- Desv. Est.: 423 caracteres
```

```
# Distribucion de lenguajes
print(f"\nDistribucion de lenguajes:")
lang_dist = df['lenguaje'].value_counts()
for lang, count in lang_dist.items():
    pct = (count / len(df)) * 100
    print(f" - {lang}: {count}, ({pct:.1f}%)")
```

```
Distribucion de lenguajes:
- c++: 848 (9.1%)
- python: 848 (9.1%)
- java: 848 (9.1%)
- javascript: 848 (9.1%)
- c#: 846 (9.1%)
- php: 846 (9.1%)
- ruby: 846 (9.1%)
- swift: 846 (9.1%)
- go: 846 (9.1%)
- kotlin: 846 (9.1%)
- fortran: 844 (9.1%)
```

```
# Balanceo de clases
vulnerable_count = (df['vulnerable'] == 1).sum()
safe_count = (df['vulnerable'] == 0).sum()
diff = abs(vulnerable_count - safe_count)

print("\nBalanceo de clases:")
print(f" - Vulnerable: {vulnerable_count},")
print(f" - Seguro: {safe_count}")
print(f" - Diferencia: {diff} registros")
if diff < 100:
    print(f" - Estado: BALANCEADO PERFECTO")
else:
    print(f" - Estado: DESBALANCEADO")

print("\n[OK] FASE 2 COMPLETADA")
```

```
Balanceo de clases:
- Vulnerable: 4,656
- Seguro: 4,656
- Diferencia: 0 registros
- Estado: BALANCEADO PERFECTO
```

[OK] FASE 2 COMPLETADA

```
# Visualizaciones
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Distribucion de clases
class_dist = df['vulnerable'].value_counts()
axes[0, 0].bar(['Seguro', 'Vulnerable'], [class_dist[0], class_dist[1]], color=['green', 'red'])
axes[0, 0].set_title('Distribucion de Clases', fontsize=12, fontweight='bold')
axes[0, 0].set_ylabel('Cantidad de muestras')

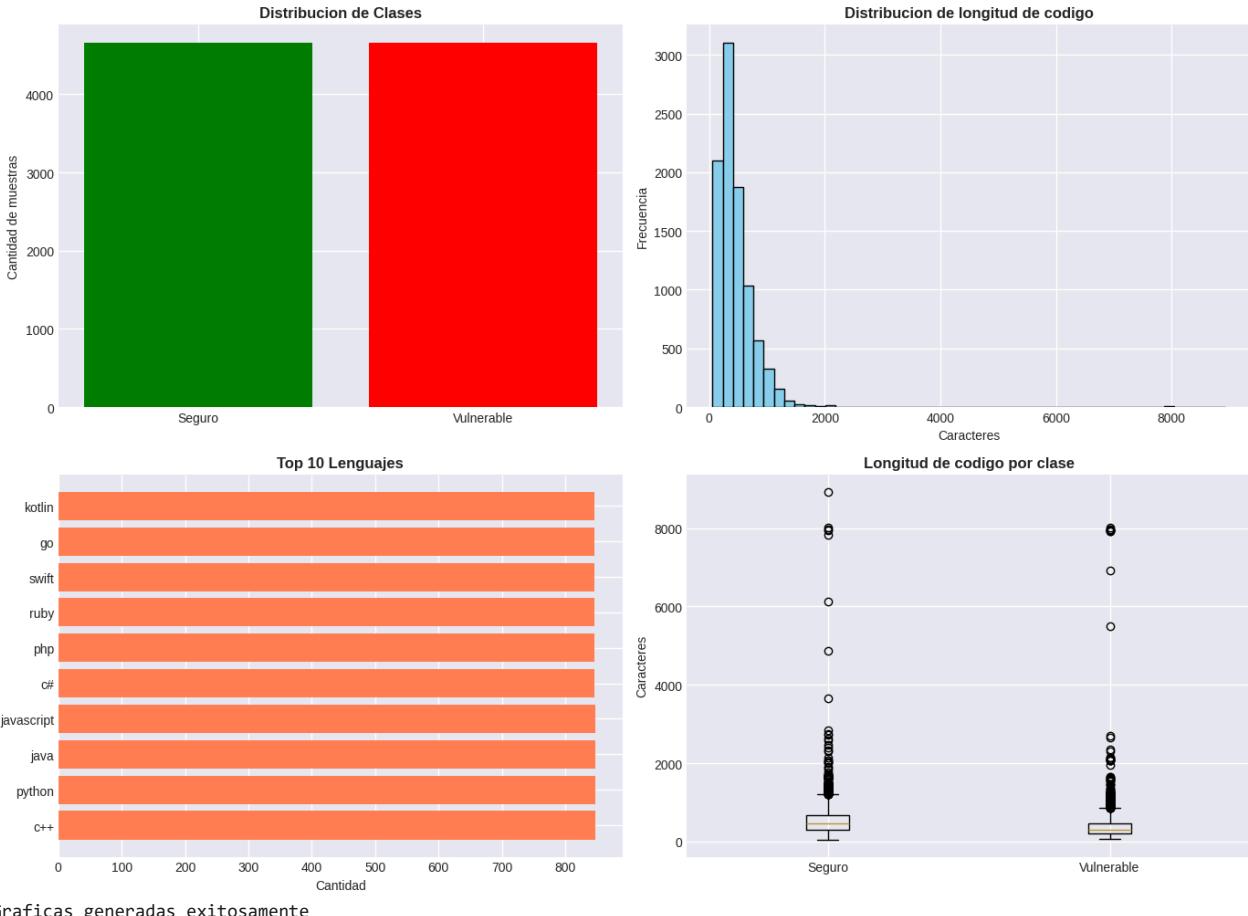
# Longitud de codigo
axes[0, 1].hist(code_lengths, bins=50, color='skyblue', edgecolor='black')
axes[0, 1].set_title('Distribucion de longitud de codigo', fontsize=12, fontweight='bold')
axes[0, 1].set_xlabel('Caracteres')
axes[0, 1].set_ylabel('Frecuencia')

# Top lenguajes
top_langs = lang_dist.head(10)
axes[1, 0].barh(top_langs.index, top_langs.values, color='coral')
axes[1, 0].set_title('Top 10 Lenguajes', fontsize=12, fontweight='bold')
axes[1, 0].set_xlabel('Cantidad')

# Longitud por clase
vulnerable_lengths = df[df['vulnerable'] == 1]['codigo'].str.len()
safe_lengths = df[df['vulnerable'] == 0]['codigo'].str.len()
axes[1, 1].boxplot([safe_lengths, vulnerable_lengths], labels=['Seguro', 'Vulnerable'])
axes[1, 1].set_title('Longitud de codigo por clase', fontsize=12, fontweight='bold')
axes[1, 1].set_ylabel('Caracteres')

plt.tight_layout()
plt.show()
```

```
print("Graficas generadas exitosamente")
```



Graficas generadas exitosamente

FASE 3: MODIFY - Feature Engineering

Objetivo: Preparar features para el modelo

```
print("\n" + "="*80)
print("FASE 3: MODIFY - Feature Engineering")
print("=*80)

# Extraer X y y
X_code = df['codigo'].values
X_lang = df['lenguaje'].values
y = df['vulnerable'].values

print(f"\nExtrayendo features...")
print(f" - Total muestras: {len(X_code)}")
print(f" - Target variable: 0=Seguro, 1=Vulnerable")
```

```
=====
FASE 3: MODIFY - Feature Engineering
=====
```

```
Extrayendo features...
- Total muestras: 9,312
- Target variable: 0=Seguro, 1=Vulnerable
```

```
# Vectorizacion TF-IDF
print("\nFeature Engineering - TF-IDF:")
print(" - Max features: 1000")
print(" - N-grams: (1, 2) - unigramas + bigramas")
print(" - Stop words: English")
print(" - Sublinear TF: True")
print(" - Min DF: 2, Max DF: 0.95")
print(" - Vectorizando codigo...")

tfidf_vectorizer = TfidfVectorizer(
    max_features=1000,
    ngram_range=(1, 2),
    lowercase=True,
    stop_words='english',
    sublinear_tf=True,
    min_df=2,
    max_df=0.95
)

X_tfidf = tfidf_vectorizer.fit_transform(X_code)
print(f"[OK] Matriz TF-IDF shape: {X_tfidf.shape}")
```

Feature Engineering - TF-IDF:
- Max features: 1000
- N-grams: (1, 2) - unigramas + bigramas
- Stop words: English
- Sublinear TF: True
- Min DF: 2, Max DF: 0.95
- Vectorizando codigo...
[OK] Matriz TF-IDF shape: (9312, 1000)

```
# Language Encoding
print("\nFeature Engineering - Language Encoding:")
lang_encoder = LabelEncoder()
X_lang_encoded = lang_encoder.fit_transform(X_lang)
print(f" - Lenguajes unicos: {len(lang_encoder.classes_)}")
print(f" - Lenguajes: {', '.join(lang_encoder.classes_)}")
print(f"[OK] Language encoding completado")
```

Feature Engineering - Language Encoding:
- Lenguajes unicos: 11
- Lenguajes: c#, c++, fortran, go, java, javascript, kotlin, php, python, ruby, swift
[OK] Language encoding completado

```
# Combinar features
X_tfidf_dense = X_tfidf.toarray()
X_combined = np.column_stack([X_tfidf_dense, X_lang_encoded])
print(f"\nFeatures combinados:")
print(f" - TF-IDF features: {X_tfidf_dense.shape[1]}")
print(f" - Language features: 1")
print(f" - Total features: {X_combined.shape[1]}")
print(f" - Total muestras: {X_combined.shape[0]}")
```

Features combinados:
- TF-IDF features: 1000
- Language features: 1
- Total features: 1001
- Total muestras: 9,312

```
# Train/Test split
print(f"\nDividiendo en train/test (80/20, stratified)...")
X_train, X_test, y_train, y_test = train_test_split(
    X_combined, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print(f"[OK] Split completado:")
print(f" - Train: {len(X_train)} muestras")
print(f" - Test: {len(X_test)} muestras")
print(f" - Train class distribution: {np.bincount(y_train)}")
print(f" - Test class distribution: {np.bincount(y_test)}")

print("\n[OK] FASE 3 COMPLETADA")
```

Dividiendo en train/test (80/20, stratified)...
[OK] Split completado:
- Train: 7,449 muestras

```
- Test: 1,863 muestras
- Train class distribution: [3724 3725]
- Test class distribution: [932 931]
```

[OK] FASE 3 COMPLETADA

▼ FASE 4: MODEL - Entrenar RandomForest

Objetivo: Entrenar modelo de detección binaria

```
print("\n" + "*80)
print("FASE 4: MODEL - Entrenar RandomForest")
print("*80)

print("\nConfiguracion del modelo:")
print(" - Algoritmo: RandomForestClassifier")
print(" - N estimators: 200 arboles")
print(" - Max depth: 25")
print(" - Min samples split: 5")
print(" - Random state: 42")
print(" - N jobs: -1 (usar todos los cores)")

print("\nEntrenando modelo...")
model = RandomForestClassifier(
    n_estimators=200,
    max_depth=25,
    min_samples_split=5,
    random_state=42,
    n_jobs=-1,
    verbose=1
)

model.fit(X_train, y_train)
print("[OK] Modelo entrenado exitosamente")

print("\n[OK] FASE 4 COMPLETADA")

=====
FASE 4: MODEL - Entrenar RandomForest
=====

Configuracion del modelo:
- Algoritmo: RandomForestClassifier
- N estimators: 200 arboles
- Max depth: 25
- Min samples split: 5
- Random state: 42
- N jobs: -1 (usar todos los cores)

Entrenando modelo...
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed:    2.1s
[OK] Modelo entrenado exitosamente

[OK] FASE 4 COMPLETADA
[Parallel(n_jobs=-1)]: Done 196 tasks      | elapsed:    9.4s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:    9.6s finished
```

▼ FASE 5: ASSESS - Evaluar performance

Objetivo: Evaluar métricas, K-Fold y análisis detallado

```
print("\n" + "*80)
print("FASE 5: ASSESS - Evaluar Performance")
print("*80)

print("\nGenerando predicciones...")
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]

# Métricas en train set
train_acc = accuracy_score(y_train, y_pred_train)
print(f"\nMETRICAS EN TRAIN SET:")
print(f" - Accuracy: {train_acc:.4f} ({train_acc*100:.2f}%)")
```

```
=====
FASE 5: ASSESS - Evaluar Performance
=====

Generando predicciones...
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 46 tasks    | elapsed:   0.1s
[Parallel(n_jobs=2)]: Done 196 tasks    | elapsed:   0.6s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.6s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 46 tasks    | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks    | elapsed:   0.1s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.1s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 46 tasks    | elapsed:   0.0s

METRICAS EN TRAIN SET:
- Accuracy: 0.8581 (85.81%)
[Parallel(n_jobs=2)]: Done 196 tasks    | elapsed:   0.1s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.1s finished
```

```
# K-Fold Cross-Validation
print("\nVALIDACION CRUZADA (5-Fold StratifiedKFold):")
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_scores = cross_val_score(model, X_combined, y, cv=kfold, scoring='accuracy', n_jobs=-1)

for i, score in enumerate(cv_scores, 1):
    print(f" - Fold {i}: {score:.4f}")
print(f" - Promedio: {cv_scores.mean():.4f} (+/- {cv_scores.std():.4f})")
```

```
VALIDACION CRUZADA (5-Fold StratifiedKFold):
- Fold 1: 0.7874
- Fold 2: 0.7944
- Fold 3: 0.7938
- Fold 4: 0.7938
- Fold 5: 0.7916
- Promedio: 0.7922 (+/- 0.0026)
```

```
# Metricas en test set
test_acc = accuracy_score(y_test, y_pred_test)
precision = precision_score(y_test, y_pred_test)
recall = recall_score(y_test, y_pred_test)
f1 = f1_score(y_test, y_pred_test)
roc_auc = roc_auc_score(y_test, y_pred_proba)

print(f"\nMETRICAS EN TEST SET:")
print(f" - Accuracy: {test_acc:.4f} ({test_acc*100:.2f}%)")
print(f" - Precision: {precision:.4f} ({precision*100:.2f}%)")
print(f" - Recall: {recall:.4f} ({recall*100:.2f}%)")
print(f" - F1-Score: {f1:.4f} ({f1*100:.2f}%)")
print(f" - ROC-AUC: {roc_auc:.4f} ({roc_auc*100:.2f}%)")
```

```
METRICAS EN TEST SET:
- Accuracy: 0.7901 (79.01%)
- Precision: 0.7373 (73.73%)
- Recall: 0.9012 (90.12%)
- F1-Score: 0.8110 (81.10%)
- ROC-AUC: 0.8883 (88.83%)
```

```
# Matriz de confusion
print("\nMATRIZ DE CONFUSION (Test):")
cm = confusion_matrix(y_test, y_pred_test)
print(f" - True Negatives (Seguro correcto): {cm[0, 0]}")
print(f" - False Positives (Falsa alarma): {cm[0, 1]}")
print(f" - False Negatives (Vulnerabilidad NO detectada): {cm[1, 0]}")
print(f" - True Positives (Vulnerable correcto): {cm[1, 1]}")
```

```
MATRIZ DE CONFUSION (Test):
- True Negatives (Seguro correcto): 633
- False Positives (Falsa alarma): 299
- False Negatives (Vulnerabilidad NO detectada): 92
- True Positives (Vulnerable correcto): 839
```

```
# Analisis de overfitting
overfitting = train_acc - test_acc
print("\nANALISIS DE OVERFITTING:")
print(f" - Train Accuracy: {train_acc:.4f}")
print(f" - Test Accuracy: {test_acc:.4f}")
print(f" - Diferencia: {overfitting:.4f} ({overfitting*100:.2f}%)")
```

```
if overfitting < 0.10:
    print(f" - Estado: ACCEPTABLE (bajo overfitting)")
else:
    print(f" - Estado: ADVERTENCIA (overfitting detectado)")
```

ANALISIS DE OVERRFITTING:

- Train Accuracy: 0.8581
- Test Accuracy: 0.7901
- Diferencia: 0.0680 (6.80%)
- Estado: ACCEPTABLE (bajo overfitting)

```
# Reporte de clasificacion
print("\nREPORTE DE CLASIFICACION:")
print(classification_report(y_test, y_pred_test, target_names=['Seguro (0)', 'Vulnerable (1)']))
```

REPORTE DE CLASIFICACION:				
	precision	recall	f1-score	support
Seguro (0)	0.87	0.68	0.76	932
Vulnerable (1)	0.74	0.90	0.81	931
accuracy			0.79	1863
macro avg	0.81	0.79	0.79	1863
weighted avg	0.81	0.79	0.79	1863

```
# Visualizaciones finales
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Matriz de confusión
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=axes[0, 0])
axes[0, 0].set_title('Matriz de Confusión', fontsize=12, fontweight='bold')
axes[0, 0].set_ylabel('Real')
axes[0, 0].set_xlabel('Predicho')

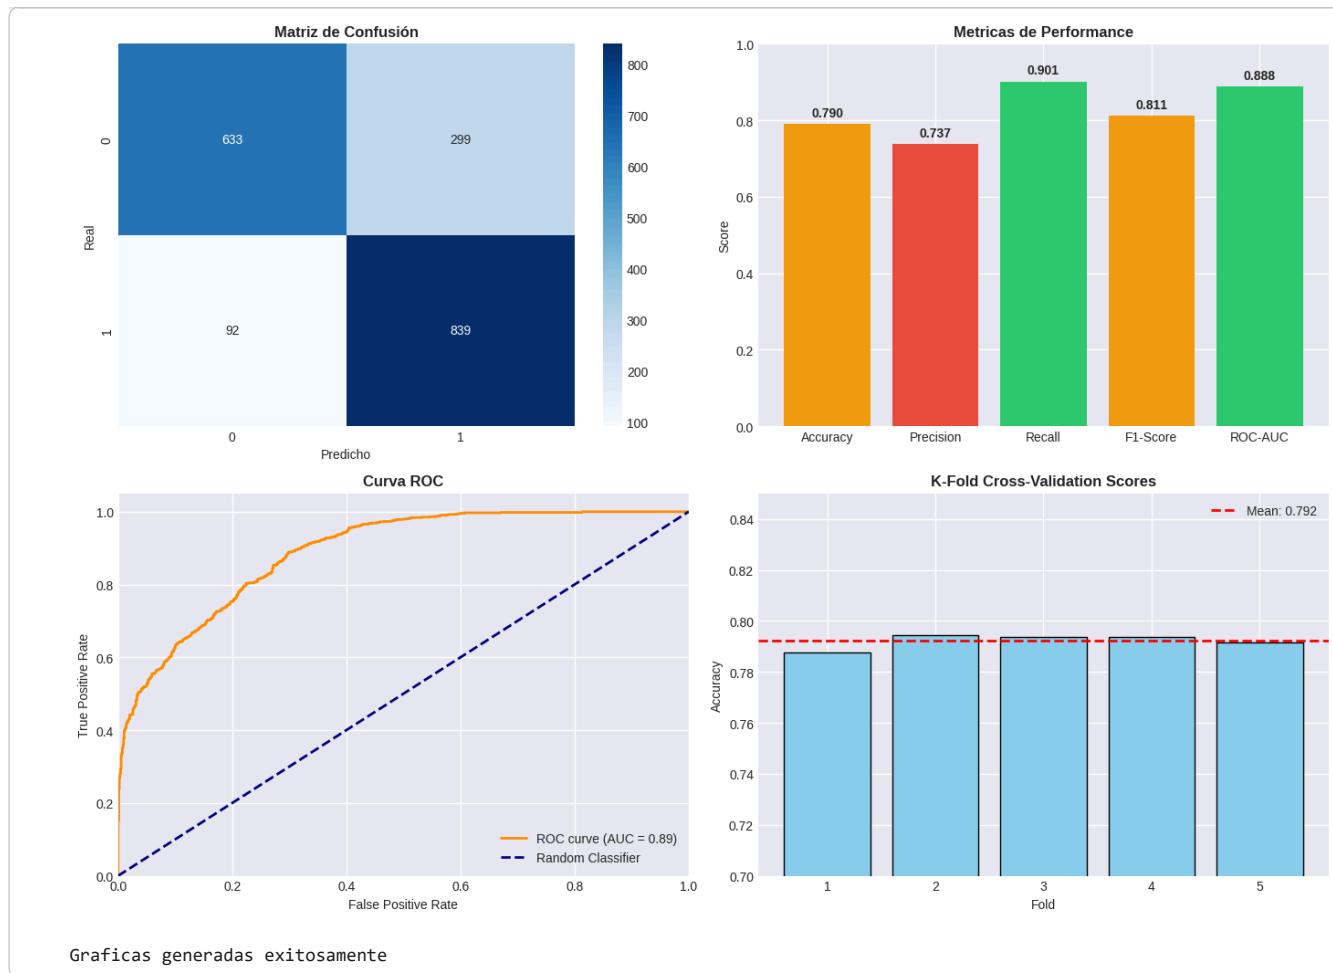
# Metricas
metricas = ['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC-AUC']
valores = [test_acc, precision, recall, f1, roc_auc]
colores = ['#2ecc71' if v > 0.85 else '#f39c12' if v > 0.75 else '#e74c3c' for v in valores]
axes[0, 1].bar(metricas, valores, color=colores)
axes[0, 1].set_title('Metricas de Performance', fontsize=12, fontweight='bold')
axes[0, 1].set_ylabel('Score')
axes[0, 1].set_ylim([0, 1])
for i, v in enumerate(valores):
    axes[0, 1].text(i, v + 0.02, f'{v:.3f}', ha='center', fontweight='bold')

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
axes[1, 0].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
axes[1, 0].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random Classifier')
axes[1, 0].set_xlim([0.0, 1.0])
axes[1, 0].set_ylim([0.0, 1.05])
axes[1, 0].set_xlabel('False Positive Rate')
axes[1, 0].set_ylabel('True Positive Rate')
axes[1, 0].set_title('Curva ROC', fontsize=12, fontweight='bold')
axes[1, 0].legend(loc="lower right")

# K-Fold scores
axes[1, 1].bar(range(1, 6), cv_scores, color='skyblue', edgecolor='black')
axes[1, 1].axhline(cv_scores.mean(), color='red', linestyle='--', linewidth=2, label=f'Mean: {cv_scores.mean():.3f}')
axes[1, 1].set_title('K-Fold Cross-Validation Scores', fontsize=12, fontweight='bold')
axes[1, 1].set_xlabel('Fold')
axes[1, 1].set_ylabel('Accuracy')
axes[1, 1].set_ylim([0.7, 0.85])
axes[1, 1].legend()

plt.tight_layout()
plt.show()

print("\nGraficas generadas exitosamente")
```



Graficas generadas exitosamente

```
print("\n[OK] FASE 5 COMPLETADA")
print("\n" + "="*80)
print("RESUMEN DEL ENTRENAMIENTO")
print("="*80)
print(f"\nPERFORMANCE FINAL:")
print(f" - Test Accuracy: {test_acc*100:.2f}%")
print(f" - Test Recall: {recall*100:.2f}% (detecta vulnerabilidades)")
print(f" - Test ROC-AUC: {roc_auc*100:.2f}%")
print(f" - 5-Fold CV: {cv_scores.mean()*100:.2f}% ± {cv_scores.std()*100:.2f}%")
print(f" - Overfitting Gap: {overfitting*100:.2f}%")
print(f"\nESTADO:  MODELO ENTRENADO Y VALIDADO")
```

GUARDAR ARTEFACTOS (OPCIONAL)

Guardar modelos para usar en producción

```
# Guardar modelos y vectorizadores
output_dir = '/content/drive/MyDrive/CVEfixes/models' # Ajusta según tu setup

try:
    os.makedirs(output_dir, exist_ok=True)

    # Guardar modelo
    with open(f'{output_dir}/vulnerability_detector.pkl', 'wb') as f:
        pickle.dump(model, f)

    # Guardar vectorizador
    with open(f'{output_dir}/vectorizer_detector.pkl', 'wb') as f:
        pickle.dump(tfidf_vectorizer, f)

    # Guardar encoder
    with open(f'{output_dir}/language_encoder.pkl', 'wb') as f:
        pickle.dump(lang_encoder, f)

    print(f"[OK] Modelos guardados en {output_dir}")
except Exception as e:
    print(f"[WARNING] No se pudieron guardar los modelos: {e}")
    print("Esto es normal si ejecutas en Colab sin Google Drive montado")
```

[OK] Modelos guardados en /content/drive/MyDrive/CVEfixes/models

- ✓ INFERENCIA - Hacer predicciones con código nuevo

```

# Ejemplo de inferencia con nuevo codigo
print("\n" + "="*80)
print("EJEMPLO DE INFERNERIA")
print("="*80)

# Ejemplos de prueba
ejemplos = [
    ("int x = 5; x++;", "c++", "Seguro"),
    ("char buffer[10]; strcpy(buffer, user_input);", "c++", "Vulnerable - Buffer Overflow"),
    ("x = 10; print(x)", "python", "Seguro"),
    ("query = 'SELECT * FROM users WHERE id = ' + user_input", "python", "Vulnerable - SQL Injection"),
]

for codigo, lenguaje, etiqueta_real in ejemplos:
    print(f"\nCodigo: {codigo[:50]}...")
    print(f"Lenguaje: {lenguaje}")
    print(f"Etiqueta real: {etiqueta_real}")

# Vectorizar
features = tfidf_vectorizer.transform([codigo])
lang_encoded = lang_encoder.transform([lenguaje])
features_combined = np.column_stack([features.toarray(), lang_encoded])

# Prediccion
pred = model.predict(features_combined)[0]
prob = model.predict_proba(features_combined)[0]

pred_label = "Vulnerable" if pred == 1 else "Seguro"
confianza = prob[pred]

print(f"Prediccion: {pred_label} (Confianza: {confianza*100:.1f}%)")
print(f"Probabilidades: Seguro={prob[0]:.3f}, Vulnerable={prob[1]:.3f}")
print("-" * 80)

Código: int x = 5; x++...
Lenguaje: c++
Etiqueta real: Seguro
Prediccion: Vulnerable (Confianza: 56.5%)
Probabilidades: Seguro=0.435, Vulnerable=0.565
-----
Código: char buffer[10]; strcpy(buffer, user_input);...
Lenguaje: c++
Etiqueta real: Vulnerable - Buffer Overflow
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.1s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.1s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.1s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.0s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.0s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.0s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.1s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.1s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:   0.0s
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:   0.1s
[Parallel(n_jobs=2)]: Done 200 out of 200 | elapsed:   0.1s finished
Prediccion: Vulnerable (Confianza: 63.5%)

```

Etiqueta real: Seguro

```
Prediccion: Vulnerable (Confianza: 54.1%)
Probabilidades: Seguro=0.459, Vulnerable=0.541
```

```
-----  
Codigo: query = 'SELECT * FROM users WHERE id = ' + user_i...  
Lenguaje: python  
Etiqueta real: Vulnerable - SQL Injection  
Prediccion: Vulnerable (Confianza: 69.8%)  
Probabilidades: Seguro=0.302, Vulnerable=0.698
```

```
-----  
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.  
[Parallel(n_jobs=2)]: Done  46 tasks      | elapsed:    0.0s  
[Parallel(n_jobs=2)]: Done 196 tasks      | elapsed:    0.1s
```